

Recovering Repeated Structure in Flat 2D Rectilinear Layouts: A Geometry-First Framework Unifying Array Detection and Hierarchy Recovery

July 2026

Contents

1. Abstract	2
2. Introduction	3
2.1 Motivation	3
2.2 The problem, and why it is not just “find the period”	3
2.3 Contributions	3
3. Related Work	4
4. Problem Statement	5
4.1 Array detection	5
4.2 Hierarchy recovery	5
4.3 Assumptions and scope	6
4.4 Non-uniqueness (stated up front)	6
5. Method	6
5.1 Scanline construction (coordinate snapping + open-slab sampling)	6
5.2 Two symbolic encodings	7
5.3 Per-line repetition and period-phase reconciliation	7
5.4 Geometric verification and defect model	8
5.5 Hierarchy recovery	9
5.6 D4 reflection and rotation	11
5.7 The unification: array detection = one-level recovery + lattice fit	11
6. Implementation	12
7. Evaluation	13
7.1 Correctness: bug-class regression tests and demos	13
7.2 The unification result	15
7.3 Hierarchy-erasure benchmark	16
7.4 Performance and scalability	16
7.5 Real-GDS hierarchy-erasure	17

8. Limitations and Non-Uniqueness	18
9. Conclusion and Future Work	19
References	20

Draft — internal working version. All numbers are from measured runs of the implemented C++/Boost.Geometry system (synthetic benchmarks and a real `gdstk` GDS round-trip). The one remaining open item — running the identical harness on a real SKY130/ORFS `.gds` crop — is marked [TBD].

1. Abstract

Flat 2D geometric layouts — integrated-circuit masks after hierarchy flattening, tiled facades, decorative patterns, printed arrays — routinely contain repeated structure that was present as *step-and-repeat* metadata upstream but is lost once the geometry is reduced to a single flat pool of polygons. Recovering that structure is valuable for compression, verification, and layout analysis, but it is hard precisely because the repetition is usually *partial*: a regular region embedded in clutter, borders, routing, and localized defects, rather than a globally periodic image.

We present a single geometry-first framework that recovers repeated structure across the full spectrum from a rigid lattice **array** to an arbitrary nested **hierarchy** of reusable cell definitions and instance placements. The framework rests on two shared ideas. First, **symbolic string / local-hash methods are candidate generators, never proofs**: they are cheap, lossy filters that propose where repetition might live; only **canonicalized geometric equality** certifies it. Second, **a minimum-description-length (MDL) objective** decides which repeated substructures are worth promoting to cells. The central claim, and the organizing thesis of this paper, is that **array detection is exactly the one-level, lattice-aligned special case of hierarchy recovery** — not merely analogous to it. We demonstrate this empirically: on the same flattened layout, the dedicated array detector and a one-level run of the hierarchy recoverer followed by a lattice fit produce **identical lattice parameters** (step vectors and grid) with equivalent total site accounting, differing only in one modal-cell attribution on the standard-cell case (§7.2).

The framework is implemented in C++17 on Boost.Geometry. It reproduces three previously-diagnosed bug classes as passing regression tests, recovers both decorative and standard-cell demo layouts (including directly-abutted zero-gap geometry that a single encoding cannot handle), recovers D4-oriented (mirrored/rotated) instances as one cell, and passes both a synthetic and a **real-GDS** hierarchy-erasure benchmark with exact flatten-equality — $3.3\text{--}3.5\times$ compression on the synthetic datapaths and $3.80\times$ on the real-GDS round-trip — recovering the leaf cell in every case. (An earlier, higher compression figure came from an over-covering verification bug; the corrected, exact-multiset verification both fixed it and lowered the honest number — verification is a correctness gate, not just a metric.) Performance is linear at roughly $3\ \mu\text{s}$ per rectangle; 1.58 M rectangles recover in about 4.8 s single-round (about 30 s under full recursion) on one core. The real-GDS evaluation writes and flattens a genuinely hierarchical layout with `gdstk` and recovers from the flat rectangles.

Throughout we follow the design documents’ insistence on honesty about scope and claims: the recovered hierarchy is *one* of many valid decompositions of the same flat geometry, so where ground

truth exists we report **agreement, not accuracy**.

2. Introduction

2.1 Motivation

A modern integrated-circuit layout is authored hierarchically: a designer defines a small cell (a logic gate, a bit-slice, an SRAM bit), and place-and-route tools instantiate it thousands or millions of times, often on regular rows and arrays with array-reference (AREF) metadata recording the step and count compactly. This hierarchy is what keeps a GDSII/OASIS file for a billion-transistor chip a manageable size. But many downstream operations — DRC region extraction, third-party IP analysis, reverse engineering of a flattened block, mask-data-prep tooling — receive geometry that has already been **flattened**: the cell and array references are expanded into raw polygons and the hierarchy metadata is gone. The same happens in non-IC settings: a decorative tiling, a rendered pattern, or a scanned facade arrives as flat geometry with its generative structure implicit.

Recovering the lost structure directly from geometry has concrete payoffs. If a region is a genuine array, reporting it as (tile, step-x, step-y, count) plus a short defect list is dramatically smaller than the expanded polygon set, and it exposes the array’s *defects* (a missing instance, an extra overlaid polygon) as first-class findings. More generally, inferring a compact hierarchy re-compresses a flattened layout and can reveal design intent that the flattening erased.

2.2 The problem, and why it is not just “find the period”

Two properties make this hard and distinguish it from classic periodicity detection:

- **Partiality.** The layout is *not* globally periodic. An array occupies a rectangular sub-region while borders, routing, filler, and unrelated blocks surround it. Multiple distinct arrays may coexist. A global-periodicity method (autocorrelation, global Fourier peak) answers the wrong question; it assumes the whole image repeats.
- **Representation slack.** The same logical geometry can be represented by different polygon decompositions — fractured versus merged, different winding, near-duplicate coordinates from floating-point noise. Any comparison that is naive about representation will reject valid repeats for non-geometric reasons.

The framework confronts both directly: repetition evidence is gathered *locally* (line-wise runs, local geometric neighborhoods) so partiality is natural, and every acceptance decision routes through **canonicalized** geometric comparison so representation slack is neutralized.

2.3 Contributions

- **A unified geometry-first framework** covering the whole range from rigid lattice arrays to arbitrary nested hierarchy, built on two shared mechanisms: canonicalized geometric verification as the sole certifier, and an MDL compression-gain objective as the promotion rule.
- **An empirical demonstration of the unification:** array detection is realized as one-level hierarchy recovery plus a lattice fit, and the two paths produce identical lattice parameters (dx, dy, grid) with equivalent total site accounting — differing only in one modal-cell attribution

on the standard-cell case (§7.2). This makes “array detection is a special case of hierarchy recovery” a demonstrated equality rather than a framing convenience.

- **Two complementary symbolic encodings** (occupancy-interval and edge-boundary) that together handle both gapped and directly-abutted (zero-gap) geometry — a case a single encoding provably fails on.
 - **A period-phase reconciliation formulation** as clustering over physical period, modular phase, and spatial support (with a center-pruning step that defeats the transitive weak-link failure mode), replacing ad hoc “try-shifts” alignment.
 - **A defect model** ($\text{missing} = \text{expected} - \text{observed}$, $\text{extra} = \text{observed} - \text{expected}$) computed with exact geometry, and correctly attributed to physical positions across defects — reproducing and guarding against three previously-diagnosed bug classes as regression tests.
 - **D4 orientation support** (the 8 axis-preserving reflections/rotations), so mirrored standard-cell rows recover as one cell rather than eight spurious ones.
 - **A linear-time implementation** and the finding that the algorithm itself was never the scaling bottleneck — the $\text{flatten}(H) = G$ verification was — with the fix that restores linearity (§7.4).
 - **An exact $O(n)$ rectangle-multiset verification used as a correctness gate**, not merely a metric: it caught an over-covering recovery bug (nested instances silently overlapping, rectangles double-counted) that a weaker $G \subseteq \text{flatten}$ check had rubber-stamped, and forced the honest compression numbers down (§6, §7.3, §7.5).
-

3. Related Work

We situate the framework against three families, briefly and honestly.

Global-periodicity methods. Autocorrelation, Fourier-peak, and lattice-from-image techniques recover a single dominant period for an image or point set. They are efficient and well-understood, but they assume global periodicity. Our target is the opposite regime — a partial array embedded in a mostly non-repeating layout — where a global period either does not exist or is dominated by clutter. We use periodicity only *line-wise and locally*, as candidate evidence, never as a global assumption.

Pairwise geometric motif matching. One can search directly for repeated geometry by comparing polygon neighborhoods against each other. This is expressive but tends toward quadratic cost in the number of primitives if done naively, and it needs a canonicalization to be robust. We keep the canonicalized geometric comparison (it is our certifier) but avoid the quadratic candidate search: array candidates come from line-wise runs (linear per line), and hierarchy candidates come from local geometric *shingles* gathered through a spatial index (an *r*tree), so candidate generation never pays the all-pairs cost.

String maximal-repetition / grammar induction. Linear-time *runs* algorithms find maximal periodic substrings; suffix structures and grammar-induction methods (Sequitur, Re-Pair) find repeated (not necessarily contiguous) substrings and are natively aligned with a compression objective. We borrow these as *serialization-channel* candidate generators, but with an explicit tool distinction the design documents make [v4, HR]: periodic *runs* fit adjacent-on-a-lattice array instances, whereas *repeated substrings* fit cell instances scattered with unrelated content between

them. Crucially, a string match over serialized geometry is lossy — serialization order is itself a choice, and tolerance binning can preserve a symbolic period while the physical lattice drifts — so the string stage is a filter and the geometry stage is the proof.

The distinguishing stance of this work is the combination: **partial** structure in clutter, **sub-quadratic** candidate generation via line-runs and local shingles, **canonicalized geometric certification**, and an **MDL** objective that ties the whole thing to description-length reduction — with array detection falling out as a special case rather than being treated as a separate problem.

4. Problem Statement

4.1 Array detection

Input. A collection of axis-aligned rectangles in the plane (the rectilinear scope means a rectangle is the only primitive we need; a fractured polygon is several rectangles).

Target. A candidate array is a tuple $\mathbf{A} = (\mathbf{B}, \mathbf{u}, \mathbf{v}, \mathbf{T})$ where $\mathbf{B}$ is a rectangular bounding box, $\mathbf{u} = (dx, 0)$ and $\mathbf{v} = (0, dy)$ are orthogonal lattice vectors, and $\mathbf{T}$ is the tile geometry of one fundamental cell. The **exact array condition** [v4, §1.2] is

$$G \cap B = \bigcup_{i,j} ((T + i\mathbf{u} + j\mathbf{v}) \cap B) \quad (\text{after coordinate snapping and canonicalization})$$

The canonicalization qualifier is load-bearing: comparing raw, unsnapped geometry would reject valid arrays for representational reasons. The detector seeks a **primitive** tile (the smallest verified repeating unit) but may report a **supercell** when primitive recovery is ambiguous; a coincidental match that survives no candidate period is a **pseudo-array** and is discarded [v4, §2.4].

Importantly, the symbolic stage never establishes the exact array condition directly. It establishes only a weaker *necessary* condition — that many horizontal slabs are periodic in $\mathbf{x}$ with period $d\mathbf{x}$, and many vertical slabs periodic in $\mathbf{y}$ with period $d\mathbf{y}$ — which geometric verification then upgrades to the full sufficient condition.

4.2 Hierarchy recovery

Input. Flat polygon geometry $\mathbf{G}$ (again, rectangles under our scope).

Output. A hierarchy $H = (C_1..C_m, I_1..I_k, R)$ of cell definitions, instance placements, and residual geometry, such that

$$\text{flatten}(H) = G \quad \text{within tolerance } \varepsilon$$

minimizing

$$\text{cost}(H) = \sum_i \text{cost}(\text{body}(C_i)) + \sum_j \text{cost}(\text{inst}(I_j)) + \text{cost}(\text{residual}(R))$$

Cells may reference smaller cells (nested hierarchy), which is how “arrays inside arrays” are represented [HR, §1.1]. Candidate generation is heuristic; only `flatten(H) = G`, checked by canonicalized equality, certifies a hierarchy.

4.3 Assumptions and scope

Shared by both problems [v4, §1.1; HR, §1.2]:

- **Axis-aligned, rectilinear (Manhattan)** geometry with an orthogonal two-vector lattice — not arbitrary oblique lattices.
- **Single-layer** geometry (multi-layer/multi-class is future work).
- **Fractured polygons** are expected; verification must tolerate one logical shape represented as several rectangles.
- **Exact repetition** (within tolerance ε) is the primary target; bounded-defect matching is a defined extension, not the default.
- Holes and overlapping polygons are out of scope for the initial version.
- **Transform scope.** Translation-only by default, deliberately **extended to the dihedral group D4** — the 8 axis-preserving orientations (rotations by 0/90/180/270°, each optionally mirrored) — because place-and-route flows routinely mirror alternating standard-cell rows to share a power rail. Under translation-only canonicalization a mirrored instance would register as unrelated geometry; instance records therefore carry an **orientation field**, not just a position [HR, §1.2].

4.4 Non-uniqueness (stated up front)

A flat layout generally admits **many** valid hierarchical decompositions. ABABABAB flattens identically whether the source was one cell $AB \times 4$, alternating A/B, or $ABAB \times 2$. The target is therefore *a compact, geometrically valid hierarchy*, not *the designer’s original hierarchy*. Where evaluation compares against a known source, we report **agreement**, not **accuracy** [HR, §1.3]. This is not a caveat bolted on at the end; it shapes the evaluation (§7) and the honest limitations (§8).

5. Method

The pipeline maps 1:1 onto the design documents’ sections and the implemented C++ modules. Array detection runs stages 5.1–5.4; hierarchy recovery reuses 5.1’s snapping and 5.4’s canonicalized-equality machinery and adds 5.5–5.6.

5.1 Scanline construction (coordinate snapping + open-slab sampling)

Scanlines are derived from the arrangement induced by rectangle vertices, with two precision rules that matter for correctness [v4, §2.1]:

- **Coordinate snapping before slab construction.** Raw coordinates are snapped to a tolerance grid (default 0.05) and coalesced, so near-duplicate coordinates do not fragment the layout into a blizzard of near-zero-height slabs. X, Y = unique snapped x-, y-coordinates.
- **Open-slab sampling.** Consecutive Y values define open horizontal slabs; each non-empty slab is represented by a sample line strictly *inside* the open interval (the midpoint), never on a boundary. For rectilinear polygons the intersection topology is constant for any horizontal

line strictly inside one open slab, so one representative line fully characterizes it. Horizontal edges lie on slab boundaries and are represented through the boundary coordinates directly.

Vertical slabs are built symmetrically by treating x as the scan-normal axis, so one code path (`build_slabs(axis_swapped)`) serves both passes. Each slab keeps two views: the **merged** covered intervals (for the occupancy encoding) and the **raw per-rectangle spans** with their heights (for the edge encoding, which must see internal boundaries the merge would dissolve). Snapping and open-interval sampling compose: snapping bounds the slab count, open sampling keeps each surviving slab unambiguous; neither alone suffices.

5.2 Two symbolic encodings

Each slab is turned into a metric-preserving token stream — every token carries both a symbol and a physical length, because tolerance binning is not automatically metric-preserving and reconciliation must operate on the physical fields [v4, §2.2].

- **Occupancy-interval encoding.** From the *merged* covered intervals: filled tokens carry the quantized covered width, gap tokens the quantized empty width. Strong when features have real gaps between them.
- **Edge-boundary encoding.** From the *raw, unmerged* per-rectangle spans, each filled token keyed on (`width`, `edge-height`), with no gap token between abutted cells. This is the concrete realization of “every edge characterized by its projection vector” specialized to axis-aligned boxes [NOTES, decision 9].

Why two encodings — the zero-gap case. Occupancy leans on gap structure. Two directly-abutted cells with zero gap merge into one wide covered interval, so occupancy cannot see the repeated boundary between them; the edge encoding keeps them as two distinct tokens and recovers the array. This is not hypothetical: the prototype implemented only occupancy and its standard-cell demo had to insert artificial gaps to work at all — a live demonstration that one encoding is insufficient. The C++ system implements both and recovers a genuinely zero-gap abutted standard-cell array through the edge encoding (§7.1). Both filled tokens get positive interned symbol ids and gap tokens negative, so downstream code reads “filled vs. empty” from the sign.

5.3 Per-line repetition and period-phase reconciliation

Per-line repetition (`analyze_row`) finds each slab’s primitive physical period. The implementation uses the prototype’s $O(n^2)$ self-consistency scan rather than a linear-time runs algorithm; both the design and the prototype note this stage is swappable without downstream impact [NOTES, decision 10]. The period chosen is the smallest one whose majority-vote canonical unit actually **reproduces** the row above a strict self-consistency floor — *not* merely the smallest that clears a loose match fraction. This is exactly what defeats bug 2 (below): a half-period that matches only because two of four widths coincide scores below the strict floor and loses to the true primitive. The loose floor survives only to flag a genuine supercell (content near-repeats at d but is truly invariant only at $k \cdot d$).

Each run is recorded with both symbolic and physical fields, and its phase is $\varphi = x_start \bmod dx$ using modular distance `dist_mod(a,b;p) = min(|a-b|, p-|a-b|)`.

Reconciliation (`cluster_runs`, `intersect_passes`) lifts the design’s key correction into code: rows within a tile need *not* have identical content — a row crossing the top of a rectangle legitimately

differs from one crossing its middle. What must agree is each row’s own internal periodicity, matched on **physical step and phase**, not on symbol content [v4, §2.4]. Two runs are compatible when

$$\begin{aligned} |dx_i - dx_j| &\leq \varepsilon_{\text{period}}, \\ \text{dist_mod}(\varphi_i, \varphi_j; dx) &\leq \varepsilon_{\text{phase}}, \\ \text{overlap}([x_{\text{start}}, x_{\text{end}}]_i, [\cdot]_j) &\geq w_{\text{min}}. \end{aligned}$$

Runs are clustered by **connected components** under this hard pairwise predicate (the graph view of the design’s clustering/voting duality). To defeat the transitive weak-link problem the design warns about (A~B and B~C compatible does not make A~C good evidence), a robust cluster center is computed — **median** period and **circular-median** phase — and any member incompatible with that center is pruned, recomputing once [NOTES, decision 2]. The circular median has no closed form (phase wraps modulo the period); it is computed as the φ minimizing $\Sigma \text{dist_mod}(\varphi, \varphi_i; p)$ searched over the observed phases plus pairwise midpoints and their antipodes [NOTES, decision 3].

The same construction runs on **vertical** slabs, producing (dy, ψ) votes. A 2D hypothesis is formed where a horizontally-supported cluster and a vertically-supported cluster overlap spatially and induce a consistent rectangle — the horizontal $\times$ vertical intersection that yields the true 2D lattice. The candidate rectangle is the **union** of member supports (the intersection would collapse the rectangle to a single row whenever member supports vary; since the compatibility predicate already requires pairwise overlap, members are genuine co-lattice evidence and the union is the honest footprint) [NOTES, decision 4]. Clusters are classified primitive / supercell / pseudo-array; LCM combines periods only for runs already classified as genuine-supercell evidence, never as a default rule.

Pseudo-code for the reconciliation core:

```
function reconcile(row_runs):
    lines = [ project(r) for r in row_runs ]          # onto (scan, normal) axes
    graph = edges { (i,j) : compatible(lines[i], lines[j]) }
    for comp in connected_components(graph):
        p_star = median(period over comp)
        phi_star = circular_median(phase over comp; period = p_star)
        comp = { r in comp : |period(r)-p_star| <= eps_p
                  and dist_mod(phase(r), phi_star; p_star) <= eps_phi }
        if |comp| >= min_members:
            emit cluster(period=p_star, phase=phi_star,
                          support=union(scan/normal extents))
    # symmetric vertical pass, then:
    return intersect(horizontal_clusters, vertical_clusters)
```

5.4 Geometric verification and defect model

Reconciliation establishes only the necessary line-wise condition; verification (`verify_candidate`) upgrades a 2D hypothesis to a certified array by canonicalized geometric comparison against the source layout [v4, §2.5]:

1. Clip the original geometry to the candidate rectangle.

2. Take the canonical tile as the **modal** cell content across all instances (not cell (0,0) blindly — so one defective reference cell cannot define the tile everyone is measured against).
3. Translate tile copies across (**dx**, **dy**) to reconstruct the rectangle.
4. Canonicalize both (snap, normalize winding, merge coincident edges, sort rings) and compare. Only matching hypotheses are accepted.

Because it is already clipping every expected instance, the same pass computes the **§6 defect map** with Boost boolean ops:

$$\begin{aligned} \text{missing} &= \text{expected} - \text{observed} && (\text{instance geometry incomplete} - \text{a structural defect}) \\ \text{extra} &= \text{observed} - \text{expected} && (\text{routing/clutter over a complete instance}) \end{aligned}$$

An instance is **clean** when **none of its own geometry is missing**; overlaid extra geometry is reported separately (**n_extra**) rather than making the instance a defect [NOTES, decision 6]. This is what lets the standard-cell demo certify the array *and* report its routing overlays and its one genuine missing cell distinctly. Critically — and this is the concrete fix for bug 3 — the defect scan walks **expected physical positions** $\mathbf{x}_{\text{start}} + \mathbf{m} \cdot \mathbf{dx}$, looking up each position’s actual geometry by physical x-window, never by token index. A missing instance is reported at its true **x**, and nothing after it desyncs [v4, §6; NOTES, decision 5].

5.5 Hierarchy recovery

Hierarchy recovery (**recover_hierarchy**) reuses 5.1’s snapping and 5.4’s canonicalized equality, and adds candidate generation, MDL scoring, selection, and recursion [HR, §2].

Channel B — geometric shingles / local hashing. Independent of any global serialization, a small local neighborhood (a seed rectangle plus its **k** nearest neighbors, **k** = 1 by default) is gathered around each rectangle through an rtree, canonicalized to a translation- and D4-invariant signature, and hashed. Neighborhoods that hash identically across locations are candidate motifs. The seed is deliberately small: a large fixed neighborhood would fold per-instance clutter into the signature and split true instances apart [NOTES, Phase-2 decision 3]. The design designates Channel B the primary, structurally sound generator because its recall does not depend on a global traversal order; the multi-serialization Channel A is a cheap but structurally incomplete supplement and is left unimplemented pending an ablation [HR, §2.3; NOTES, Phase-2 decision 2].

Seed-and-extend growth. Each small seed is grown geometrically: for each occurrence, add a neighboring primitive only if it is present at a **consistent offset around every occurrence simultaneously**. Growth stops when no further addition holds across all instances, and is rtree-local and size-capped (**max_cell_members**) so a perfectly periodic region cannot grow an unbounded body. Because clutter is not present at a consistent offset around every instance, it cannot attach — this is what keeps clutter out of cell bodies [HR, §2.4].

MDL compression gain. For a candidate body **M** occurring **k** times,

$$\text{gain}(M) = (k - 1) \cdot |\text{body}(M)| - k \cdot \text{cost_instance} - \text{def_overhead}$$

The instance term scales with **k** (each occurrence needs its own reference); token/edge counts stand in for body cost, a first-pass model the design explicitly sanctions [HR, §2.6]. Only positive-gain candidates are promotable.

Greedy set-packing selection. Overlapping and nested candidates are expected (ABCD, BCD, ABCDABCD can all be valid). Selection is weighted set-packing — maximize $\sum \text{gain}(C_i)$ subject to selected covers being disjoint or properly nested. This is *not* weighted interval scheduling (covers are arbitrary 2D primitive sets with no total order, so the interval DP does not transfer). We use the greedy “sort by gain, take the best, remove conflicts, repeat” heuristic — the design’s sanctioned first pass — with ILP/global optimization left as future work [HR, §2.7; NOTES, Phase-2 decision 5].

An alternative promotion criterion — the repeat-length drop-off. MDL gain is not the only way to fix the cell scale. The design also proposes a *repeat-length drop-off* [HR, §3]: as a motif is grown, count the distinct motifs of each size L that still recur; the count stays low while the growing neighborhood is the shared cell, then rises sharply once it reaches into non-repeating surroundings — the size at the top of that rise is the natural cell scale (“grow the motif until the repeat suddenly stops”). We implement this both as a first-class inspectable analysis (`dropoff_curve`, giving $R(L)/F(L)$) and as a selectable promotion criterion (`Selection::DropOff`) that promotes motifs at the drop-off scale rather than by gain, with `Selection::MDLGain` remaining the default. Consistent with the design’s framing of §3 as a *scale prior* rather than a universal decider, the drop-off is a clean cell-scale signal for **isolated** motifs surrounded by non-repeating content — it recovers the 2-rectangle nested motif and the 4-rectangle defective motif exactly — but weakens on a **dense lattice**, where sub-units repeat as well and the sharp cliff shifts to the array extent; there, MDL gain plus exact verification remain stronger. On the clean demos the two criteria agree; exposing both makes them comparable per layout.

Recursive nesting. After replacing selected instances with cell references, the compressed representation (references + residual) is re-tokenized and passed back through the same steps, discovering cells built from cells — a datapath built from bit-slices built from gates. Non-recursive passes under-recover structure that only becomes visible after lower-level repetition is factored out [HR, §2.8]. The recursion, greedy tie-breaks, and knob settings jointly determine *which* member of the valid-hierarchy family is returned; this is the §4.4 non-uniqueness compounding across levels, and is expected, not a bug.

Exact rectangle-multiset certification — the correctness gate. This is the central systems contribution, not an implementation detail: it is the check that makes every recovered hierarchy trustworthy, and it is what caught the over-covering bug discussed below. Every instance’s membership is an exact index match on (`cell`, `orient`, `snapped-anchor`, `shape`), and the whole result is certified by an **exact per-rectangle multiset** check: every real rectangle of G must be reproduced by `flatten(H)` with exactly its multiplicity. This is strictly stronger than a $G \subseteq \text{flatten}(H)$ containment test, and the difference is load-bearing. Under containment, *over-covering* passes silently — nested instances can overlap and reproduce a rectangle twice — so containment will rubber-stamp an invalid tiling; the multiset check rejects it (under-covering shows up as missing geometry, over-covering as an invalid double-count). Two mechanisms keep recovery on the valid side of this gate: growth rejects any member that geometrically overlaps the body and maintains a per-instance *claimed-primitive* set so each primitive is claimed once (no two cell instances share a rectangle), and recovery **validates after every round and reverts any round that breaks exactness** (recurse-only-while-valid) — so isolated nesting recurses fully while a dense periodic region stops early rather than producing an invalid deeper hierarchy. The check remains $O(n)$ (§6). For axis-aligned rectangles under D4 this orientation-aware exact match *is* the canonicalized-equality check, so a separate per-candidate Boost equality pass is redundant. The one case not yet covered is **fractured polygons** (one logical shape as several rectangles that match only after edge-

merging); reusing the array verifier’s canonicalization before comparison would be needed there, and is flagged [NOTES, Phase-2 decision 6; NOTES, “Real-GDS hierarchy-erasure”].




Figure 1: Why exact rectangle-multiset certification matters. Greedy growth over a dense periodic region can place two cell instances that overlap, reproducing a rectangle twice. A containment check ($G \subseteq \text{flatten}(H)$) accepts this invalid tiling; the exact multiset check requires every real rectangle to appear with exactly its multiplicity and rejects the double-count. This is the check that caught the over-covering bug.

Defect-tolerant recovery. After clean cells are established, a rescan accepts **partial** instances (a strong majority of members present, `defect_min_support`, capped by `defect_max_missing`) and records absent leaf members as a defect layer. Under the multiset gate the constraint is that every *real* rectangle of G is still reproduced exactly; the idealized-but-absent members are the defect layer ($\text{flatten}(H) - G$, computed by Boost), so clutter overlapping a missing member is handled (that overlap is exactly what breaks a naive wholesale-rectangle subtraction). Missing *sub-cell references* are never tolerated, so nesting stays exact and clean tests are untouched [NOTES, “Array detection = the lattice-aligned special case”; NOTES, “Real-GDS hierarchy-erasure”].

5.6 D4 reflection and rotation

For rectilinear geometry the only rigid transforms keeping rectangles axis-aligned are the 8 elements of D4 (rotations 0/90/180/270°, each optionally mirrored) — exactly the GDS orientation set. The `d4.hpp` module encodes these as signed permutation matrices with group composition. Grouping is by a **D4-canonical signature** (the minimum over all 8 orientations of a neighborhood’s signature); matching tries all 8; growth and flatten **compose orientations through nesting** (a sub-cell reference’s own orientation composes with its parent’s). Every instance carries an **orient** field. This is what turns eight mirrored/rotated copies of a motif into **one** cell definition with eight oriented instances, rather than eight unrelated cells — the mirrored-standard-cell-row case made to work on purpose [HR, §1.2/§5.2; NOTES, Phase-2 decision 1].

5.7 The unification: array detection = one-level recovery + lattice fit

The two problems are one. A cell whose recovered instances happen to lie on a uniform 2D lattice *is* an array. Concretely, run hierarchy recovery constrained to **one level**, take the recovered cell’s instance anchors, and fit a regular axis-aligned lattice (`fit_lattice`) to them. The resulting (`dx`, `dy`, `ncols`, `nrows`, `occupied`, `missing`) is precisely what the dedicated array detector reports directly. §7.2 shows this holds exactly on real demo data.

Two honest wrinkles surface here, both informative. First, **unconstrained** recovery need not return the array view: it can pick a different valid decomposition (on the standard-cell layout it returns two cells at one level rather than the single four-rectangle array tile — §7.2), so “which hierarchy” is a knob (`max_levels`, tie-breaks), not a fact — the §4.4 non-uniqueness, live. Second,




Figure 2: The unification. Both algorithms consume the same flat geometry. The array detector yields a tile, step vectors, and a defect map directly; hierarchy recovery yields cells, instances, and residual, and a *one-level* run followed by a lattice fit on the recovered instances reproduces the same lattice parameters (dashed). Array detection is the one-level, lattice-aligned special case (§5.7, §7.2).

the correctness gate (§5.5) shapes what unconstrained recovery *can* do here: greedy MDL is tempted to keep pairing instances into a deeper binary-doubling tree over a dense array, but any round that would make instances overlap breaks the exact-multiset check and is reverted, so on these dense arrays unconstrained recovery correctly stops shallow rather than manufacturing an invalid deeper hierarchy. Achieving *valid* deep compression of a dense lattice needs an explicit array node, not greedy doubling (§8).

6. Implementation

The system is C++17 on **Boost.Geometry** (header-only, permissive license, chosen over CGAL because the scope is rectilinear and CGAL’s exact-kernel **Arrangement_2** machinery is more than this problem needs). Boost provides the box/polygon primitives, the boolean set operations for the defect computation and canonicalized equality, and `boost::geometry::index::rtree` for the local-neighborhood queries hierarchy recovery needs. The visualizer emits SVG directly with no extra dependency.

Module layout (each maps to a design section):

File	Stage
<code>slab.{hpp,cpp}</code>	§5.1 scanlines: snapping + open-slab sampling
<code>encoding.{hpp,cpp}</code>	§5.2 both encodings (occupancy merged; edge raw/height-keyed)
<code>runs.{hpp,cpp}</code>	§5.3 per-line repetition + §5.4 defect-tolerant scan
<code>reconcile.{hpp,cpp}</code>	§5.3 compatibility predicate, CC clustering, median/circular-median, $H \times V$ intersection
<code>verify.{hpp,cpp}</code>	§5.4 canonicalized verification + missing/extra defect map
<code>detector.{hpp,cpp}</code>	array driver + dedup + candidate-reduction funnel counts
<code>hierarchy.hpp</code>	H = cell defs + placements + residual + defect layer

File	Stage
<code>recover.{hpp,cpp}</code>	§5.5 shingles → grow → MDL → greedy → recurse; <code>flatten(H)=G</code>
<code>d4.hpp</code>	§5.6 the 8 axis-preserving orientations
<code>bench.{hpp,cpp}</code>	§7.3 hierarchy-erasure benchmark generator + scorer
<code>svg / hr_svg</code>	visualization

Performance-relevant implementation choices (§7.4 quantifies them): rtree neighbor queries for shingle k-NN and growth candidates (replacing an $O(n^2)$ per-seed distance scan); an **integer-keyed spatial hash** (FNV over `(cell, orient, snapped-anchor)`) instead of `std::string` keys in the hot lookup path; rtree-local, size-capped growth; and an **$O(n)$ exact rectangle-multiset** verification that groups and reconstructs existing rectangles exactly (never merging or refracturing) and compares multiplicities by hashing snapped rectangles, instead of the iterated Boost `union_` that made the first correctness-focused pass $O(n^2)$.

Verification is a correctness gate, not just a metric — and a linear one. The $O(n)$ multiset check is what makes recovery trustworthy at scale, and it is also what *caught* the over-covering bugs. An earlier version had relaxed the check to $G \subseteq \text{flatten}(H)$ (containment), which is cheap but too weak: it accepts a hierarchy in which two nested instances overlap and reproduce a rectangle twice. That masked a real growth bug — in a perfectly periodic region, seed-and-extend would add a neighbor’s rectangle “present at all instances,” making adjacent cell instances overlap — and inflated the reported compression, because the invalid, over-covering hierarchy looked smaller than any valid one. Switching to the exact multiset check (every real rectangle reproduced with exactly its multiplicity; over-covering rejected as an invalid double-count; flatten-only rectangles retained as the genuine defect layer) surfaced the bug, and combined with the growth fix (overlap rejection + per-instance claimed-primitive sets) and recurse-only-while-valid, it both restored correctness and *lowered* the honest compression numbers (§7.3, §7.5). The check stays $O(n)$: grouping and multiset comparison are hashing passes over snapped rectangles, no Boost boolean union in the hot path.

7. Evaluation

All numbers below are from actual runs of the implemented system, including the §7.5 real-GDS round-trip; the only open item is the real SKY130/ORFS crop, marked [TBD].

7.1 Correctness: bug-class regression tests and demos

The three bug classes diagnosed in the Python prototype are reproduced as passing regression tests (`test_bugs`), plus the closed encoding gap:

- **Bug 1 (symbol collision).** A layout where non-repeating clutter has, by coincidence, the exact width of a real repeated token. The detector still finds $dx = 20$ (the true period), the 8×6 grid, and excludes the colliding clutter from the region; all 48 instances are clean and the clutter is not counted. Non-periodic content is trimmed against the canonical unit and geometric verification is mandatory, so a single-token width coincidence can neither form a period nor survive verification.
- **Bug 2 (spurious sub-period from equal widths).** A repeating unit with an internal repeated sub-value. The detector finds $dx \approx 78$ (the full primitive period), *not* the ≈ 39

half-period, because the period is chosen by self-consistency relative to the best achievable, not an absolute floor.

- **Bug 3 (phase attribution across a defect).** A row/grid with a genuine missing instance in the middle. Of 48 expected instances, exactly 47 are clean and exactly one defect (type MISSING) is reported, **attributed to column 3, row 2 — not shifted right**. The direct row-scan test confirms the mechanism: period $dx = 20$, 8 expected positions, one defect at physical $x = 60$ (start of instance 3), not shifted. The defect scan walks expected physical positions, never token indices.
- **Zero-gap gap closure.** A directly-abutted standard-cell array (which occupancy alone cannot see) is detected via the edge encoding.

Array-detection demos (**demo**), with the candidate-reduction funnel N symbolic runs $\gg M$ clusters $\gg K$ verified arrays:

Demo	Rects	Runs N	Clusters M	Verified K	dx	dy	Grid	Clean / Missing / Extra
bauhaus (decora- tive)	101	58	2	1	40	40	8×6	47 / 1 / 0
standard- cell (gapped rows)	109	109	5	1	87	26	5×5	24 / 1 / 15
standard- cell (zero-gap abutted)	102	45	2	1	79	26	5×5	24 / 1 / 0

The gapped standard-cell demo carries 15 routing overlays (**extra**) and one genuine missing cell (**missing**, area 356.4), reported distinctly — the value of the defect model. The abutted demo isolates the encoding point: edge encoding recovers zero-gap geometry occupancy cannot. (Per NOTES, the abutted demo deliberately uses non-crossing border clutter, because edge-only encoding is clutter-fragile over zero-gap geometry — see §8.)

Hierarchy-recovery demos (**hr_demo**):

Demo	Leaf rects	Cells	Levels	Result
nested motif (irregular placement)	29	2	2	leaf cell $\times 13$, super-cell (3 refs to leaf) $\times 3$, 7 top, 3 residual; $1.93 \times$
oriented motif (all 8 D4)	18	1	1	one cell, 8 instances spanning 00..07; $1.50 \times$
defective motif	20	1	1	one 4-member cell $\times 5$, one instance missing a member; defect area 120

All three satisfy $G \subseteq \text{flatten}(H)$. The oriented case is the D4 payoff: eight mirrored/rotated copies collapse to one cell. The defective case exercises defect-tolerant recovery end to end.




Figure 3: Array detection on the decorative *bauhaus* demo. The detected 8×6 array region (blue outline) excludes the non-repeating decorative border, title bar, and corner ornaments; the dashed lattice shows the recovered $dx = dy = 40$ step; the red $\times$ marks the single deliberately-missing motif. Result: 47 clean instances, 1 defect, exact.

7.2 The unification result

On identical flat input, the dedicated array detector and one-level hierarchy recovery + lattice fit (`unify_demo`, `test_unify`) produce identical lattices:

Layout	Path	dx	dy	Grid	Occupied	Missing	flatten=G
bauhaus	array	40	40	8×6	47	1	—
(101 rects)	detector						
bauhaus	1-level	40	40	8×6	47	1	yes
	recovery +						
	lattice fit						
standard-	array	87	26	5×5	24	1	—
cell (109	detector						
rects)							
standard-	1-level	87	26	5×5	23	2	yes
cell	recovery +						
	lattice fit						

`test_unify` asserts these match: array dx/dy /grid equal the lattice-fit values, and occupied + missing equals the array instance count. (On standard-cell the two paths partition the same sites slightly differently between “occupied” and “missing” — $23+2$ vs $24+1$ — because one-level recovery groups by exact modal-cell membership while the array verifier’s modal-tile clip counts one borderline instance as clean; both agree on the lattice and both satisfy their correctness constraint.) The unconstrained full-hierarchy path finds a *different* valid decomposition on each layout — bauhaus: 1 cell at 1 level, cost 56 vs. flat 101 ($1.80\times$); standard-cell: 2 cells at 1 level, cost 43

vs. 109 ($2.53\times$) — a live instance of §4.4 non-uniqueness. (Under the corrected exact-multiset gate, unconstrained recovery on these dense arrays now stays shallow rather than manufacturing the deeper binary-doubling tree an earlier over-covering path reported; see §5.5, §5.7.) **This equality between the two independent paths is the paper’s core empirical claim: array detection is the one-level lattice-aligned special case of hierarchy recovery, demonstrated, not asserted.**

7.3 Hierarchy-erasure benchmark

The benchmark (`bench`, `erasure_demo`) builds a design with **known** hierarchy — a leaf 4-rectangle “gate” → a horizontal array “slice” → a vertical stack of alternately-mirrored slices “block” → a row of blocks “chip” — flattens it to raw rectangles, strips all cell/instance/array metadata, and asks the recoverer to reconstruct a compact hierarchy from geometry alone. Ground truth is the design’s own construction, so nothing is hand-labeled. Compression is `flat leaf count / hierarchy cost`.

Case	Rects	Cells	Levels	Top insts	Residual	Defective	Compression	flatten=G (exact)	Leaf found
flat array (1 block)	96	1	1	24	0	0	3.43 $\times$	YES	found
nested $\times 2$ blocks	196	1	1	47	8	0	3.32 $\times$	YES	found
nested + mirrored rows	246	1	1	59	10	0	3.37 $\times$	YES	found
larger + mirrored + clutter	588	1	1	142	20	0	3.54 $\times$	YES	found

Every case satisfies the exact rectangle-multiset flatten check (not merely containment) and recovers the ground-truth leaf gate; compression is $\sim$ **3.3–3.5** $\times$. The mirrored cases exercise D4 (adjacent rows share a rail). These numbers replace an earlier, higher figure ($\sim$ 5–8 $\times$) that came from the over-covering verification bug described in §6: under the weak $G \subseteq \text{flatten}$ check, invalid hierarchies whose nested instances overlapped were accepted and looked artificially smaller. The corrected exact-multiset gate both rejects those and reports the honest, valid compression. Note also that the recovered hierarchies now stay at **one level** on these dense arrays: recurse-only-while-valid reverts any deeper round that would break exactness, so greedy binary doubling cannot fragment the lattice into an invalid tree. The consequence is that *valid deep compression of a dense array is left on the table* — capturing it needs an explicit array node (represent a K-instance lattice as one array node rather than a log-K doubling tree), which is future work (§8). Defect tolerance is also demonstrated in recovery: on a defective datapath the missing gate member is recorded as the defect layer while every real rectangle is still reproduced exactly.

7.4 Performance and scalability

The recovery pipeline was profiled after a first correctness-focused pass revealed an $O(n^2)$ blowup (a $\sim$ 1000-rectangle design took minutes). Phase-by-phase timing showed **the algorithm was never the bottleneck**: the per-round work (shingle / group / build / dedup / greedy) was already linear and fast — about **6 ms at $n = 3072$** . The entire $O(n^2)$ lived in the final `flatten(H) = G` verification, which built one giant multipolygon by iterated `Boost union_`.

The four fixes (rtree neighbor queries, integer-keyed spatial hash, rtree-local size-capped growth, and the $O(n)$ exact rectangle-multiset verification — §6) removed the quadratic term entirely. The resulting scaling:

Metric	Value
Per-rectangle cost	$\sim 3 \mu\text{s}/\text{rect}$, linear
Per-round work at $n = 3072$	$\sim 6 \text{ ms}$
1.58 M rectangles, single round	$\sim 4.8 \text{ s}$ (one core)
1.58 M rectangles, full 8-level recursion	$\sim 30 \text{ s}$ (one core)
Correctness (exact rectangle-multiset flatten check)	intact at all sizes

The key methodological finding: **the interesting bottleneck was verification, not detection**. Once verification is made linear, the whole pipeline is linear — and, as §6 recounts, making that verification *exact* (rather than a cheaper containment test) is also what caught the over-covering bug, so the correctness gate and the performance fix are the same piece of code.

Path to billions (architecture ready, not yet built): **tiling + parallelism** (rounds, tiles, and candidates are independent; the rtree and per-tile recovery parallelize cleanly, with a stitch pass for motifs crossing tile borders); **streaming / out-of-core** for layouts exceeding RAM; **array-node compression** (represent a large lattice as one $O(1)$ array node, which also captures the valid deep compression §7.3 currently leaves on the table); and applying the same $O(n)$ rectangle approach to the array detector’s own verifier, which still uses the iterated-union pattern.

7.5 Real-GDS hierarchy-erasure

The real-GDS round-trip benchmark exercises genuine GDS I/O and real flattening (the geometry is synthetic; a real open-source GDS crop is the remaining step, below). `scripts/gds_roundtrip.py` uses `gdstk` to **write** a genuinely hierarchical `.gds`: a leaf **gate** cell (4 rectangles), a **slice** built as an array reference (AREF) of the gate, a **block** of alternately `x_reflection`-mirrored slices (the way place-and-route mirrors standard-cell rows to share a power rail), a **top** of blocks, plus routing clutter. It then **reads the file back and flattens it with the tool**, discarding all cell/instance/array metadata, and exports the flat rectangles. `build/gds_recover` recovers a hierarchy from those flat rectangles alone. The geometry is synthetic, but the GDS format, the hierarchy, the mirroring, and the flatten are all real — and dropping in a real **SKY130/ORFS** `.gds` uses the identical read $\rightarrow$ flatten $\rightarrow$ export path.

Measured result:

Quantity	Value
Flattened rectangles (G)	772
Ground truth	4 designer cells (gate/slice/block/top), 4 levels, gate = 4 rects $\times$ 192 instances
Recovered	1 cell of 4 rectangles , 1 level, 191 top placements, 8 residual rectangles
Leaf gate recovered	YES (4-rectangle cell)
<code>flatten(H) = G</code>	YES, exact (rectangle-multiset check)
Compression	3.80$\times$ (flat 772 $\rightarrow$ hierarchy 203)

The recoverer reconstructs the ground-truth 4-rectangle gate cell at **191 of 192** instances; one edge gate falls to residual, and the clutter is left as residual — exactly the kind of principled dis-

gds_erasure (1 cells, 1 levels, flatten == G)




Figure 4: Recovery from a flattened, metadata-stripped real GDS. The hierarchical `.gds` (gate $\rightarrow$ slice $\rightarrow$ block $\rightarrow$ top, with mirrored rows) is written and flattened by `gdstk`; the recoverer reconstructs the 4-rectangle gate cell as instances across the layout (blue), with routing clutter left as residual (gray), `flatten(H) = G` exact.

agreement §4.4 predicts. We therefore report this as **agreement, not accuracy** [HR, §1.3]: the recovered hierarchy is a compact, geometrically exact decomposition (`flatten(H) = G` holds exactly under the multiset gate), which need not match the designer’s four-level gate/slice/block/top decomposition. The $3.80\times$ compression sits at the top of the corrected $\sim 3.3\text{--}3.8\times$ range and, like the synthetic cases, is one-level — the deeper valid compression of the dense mirrored array again awaits array-node support (§8).

This is also the evaluation that surfaced the over-covering verification bug (§6): the dense, mirrored, abutting arrays in a real flattened GDS are precisely the stress case that made adjacent-instance overlap possible, which the exact-multiset gate then rejected.

The remaining source per the design’s three-source scoping recommendation [v4, §7.3] — a real flattened crop from the **SkyWater SKY130 PDK / OpenROAD-flow-scripts (ORFS)** flow — plugs into the same path; **FreePDK45**/Nangate-style libraries are a classic academic baseline and **ALIGN** analog layouts (matched-device / capacitor arrays) a harder stretch target. [TBD] run the identical harness on a real SKY130/ORFS `.gds` crop and report the same metrics row.

8. Limitations and Non-Uniqueness

We follow the design documents in being explicit about what the framework does *not* claim.

Recovered hierarchy is one of many valid decompositions. This is the central honesty point [HR, §1.3, §2.8]. A flat layout admits many geometrically valid hierarchies; both the greedy selection (order-dependent) and the recursion (operating on whatever partial compression the previous level produced) mean the output is one member of a family, not a canonical answer. We saw it live in §7.2 (unconstrained recovery picks a different decomposition from the array view) and in §7.5 (191 of 192 gates recovered, the 192nd falling to residual — a principled disagreement, not an error). Where a source hierarchy is known, we therefore report **agreement, not accuracy** — a well-formed inferred hierarchy can legitimately disagree with a designer’s decomposition while flattening to identical geometry. The knobs (`max_levels`, `min_cell_members`, `cost_instance`, `gain_min`, ...) are the controls that select *which* valid hierarchy is returned; they are non-uniqueness controls, not accuracy dials.

Greedy selection is non-optimal. Selection is weighted set-packing (NP-hard in general); the greedy heuristic has no optimality guarantee. ILP for small/medium candidate counts, and a principled global objective, are explicit future work [HR, §2.7].

Edge-only encoding is clutter-fragile. Occupancy leans on gap structure to stay robust to routing wires crossing an array (the gapped standard-cell demo has crossing wires and detects perfectly). The edge encoding has no gaps to anchor on, so clutter crossing **zero-gap** geometry perturbs the token stream and can push individual rows to spurious periods. The abutted demo therefore uses non-crossing border clutter, isolating the point it exists to make (edge encoding recovers directly-abutted geometry occupancy cannot) without implying an edge-encoding clutter-robustness the design does not yet claim. An occupancy-vs-edge ablation over routed real layouts is open [NOTES, “One honest limitation”; v4, §9 Q9].

Fractured-polygon verification not yet implemented. Recovery’s exact rectangle-multiset match certifies one logical shape only when its rectangles match without edge-merging. A shape fractured differently between two instances would need the array verifier’s canonicalization (merge coincident edges before comparison) reused in the recovery path — flagged, not built [NOTES, Phase-2 decision 6].

Valid deep compression of dense arrays is left on the table. Greedy MDL is tempted to compress a dense lattice with a deep binary-doubling tree, but on a contiguous array such doubling makes nested instances overlap, which the exact-multiset gate (§6) correctly rejects — so recovery reverts those rounds and stops at one level. This is the honest post-fix behavior: every benchmark case (§7.3) and the real GDS (§7.5) recovers the leaf and flattens exactly, but at a modest $\sim 3.3\text{--}3.8\times$ rather than the deeper compression a dense array could in principle admit. (An earlier, higher figure of $\sim 5\text{--}8\times$ was the *invalid* over-covering path that the weak containment check accepted; it is not a target to recover.) Capturing valid deep compression needs **array-node compression** — represent a K-instance lattice as one explicit array node ((`tile`, `dx`, `dy`, `count`)) rather than a placement list or a log-K doubling tree — which is also on the scalability path (§7.4). This is exactly where the framework closes back on itself: an array node *is* the array detector’s output, so array-node compression is the natural mechanism for a recovered cell whose instances happen to lie on a lattice. Future work.

Other scoped-out items. Multi-layer / multi-class geometry, holes and overlapping polygons, oblique (non-orthogonal) lattices, and the §4 candidate-scoring formula (normalized-product vs. log-weighted-sum, left unimplemented because verification alone suffices at demo scale [NOTES, decision 8]) are all out of the current scope.

9. Conclusion and Future Work

We have presented a single geometry-first framework for recovering repeated structure in flat 2D rectilinear layouts, spanning rigid lattice arrays through arbitrary nested hierarchy, unified by two shared mechanisms: canonicalized geometric equality as the only certifier (symbolic string and local-hash methods are candidate generators, never proofs), and an MDL compression-gain objective as the promotion rule. The unification is not rhetorical — array detection is realized as one-level hierarchy recovery plus a lattice fit, and the two independent paths produce identical lattice parameters (step vectors and grid) with equivalent total site accounting on real demo data, differing only in one modal-cell attribution on the standard-cell case (§7.2). The system reproduces three

previously-diagnosed bug classes as passing regression tests, handles zero-gap abutted geometry through a second encoding, recovers D4-oriented instances as one cell, and passes both a synthetic and a real-GDS hierarchy-erasure benchmark with exact rectangle-multiset flatten-equality and $\sim 3.3\text{--}3.8\times$ compression, recovering the leaf cell in every case, running linearly at $\sim 3\ \mu\text{s}/\text{rect}$ (1.58 M rectangles in ~ 4.8 s single-round on one core). A methodological point worth carrying forward: the exact verification is a correctness gate, not a scoreboard — it caught an over-covering bug that a cheaper containment check had hidden, and lowering the reported compression to the honest range was the *right* outcome, not a regression. The framework is honest about non-uniqueness throughout: agreement, not accuracy.

Future work, in rough priority order: (1) run the identical **gdstk** harness on a real SKY130/ORFS .gds crop (§7.5); (2) **array-node compression** to represent dense lattices as $O(1)$ array nodes — capturing the valid deep compression currently left on the table and closing the framework on itself, since an array node is exactly the array detector’s output; (3) **tiling + parallelism** and streaming toward billion-polygon layouts; (4) **fractured-polygon verification** in the recovery path via reused canonicalization; (5) the **Channel A** multi-serialization generator and an occupancy-vs-edge and Channel-A-vs-B **ablation**; (6) ILP/global selection to replace greedy; and (7) multi-layer support. Papers on either algorithm should follow real, benchmarked output — this draft is the working record toward that, not a finished result.

References

- [v4] *Geometry-First Array Detection in 2D Polygon Layouts via 1D Repetition Hypotheses*, v4 ([repeat-cell-detection-proposal-v4.md](#)). The array-detection design of record.
- [HR] *Hierarchy Recovery from Flattened Polygon Layouts via Repetition-Driven Compression*, v1 ([hierarchy-recovery-proposal-v1.md](#)). The companion hierarchy-recovery design.
- [NOTES] *Phase 1 — implementation notes & flagged design decisions* (NOTES.md). Design decisions taken where the proposals left choices open, measured performance, and honest limitations.
- SkyWater **SKY130** PDK: <https://skywater-pdk.readthedocs.io/en/main/contents.html>
- **OpenROAD-flow-scripts (ORFS)**: <https://openroad-flow-scripts.readthedocs.io/en/latest/>
- OpenROAD Project (background): https://en.wikipedia.org/wiki/OpenROAD_Project
- **ALIGN** (analog layout automation): <https://align-analoglayout.github.io/ALIGN-public/>
- **FreePDK45** / **Nangate**-style open standard-cell libraries (classic academic baseline).
- **DeepPCB** (secondary raster baseline only): <https://github.com/tangsanli5201/DeepPCB>